

WEB DATA

REST, Serialization, Views

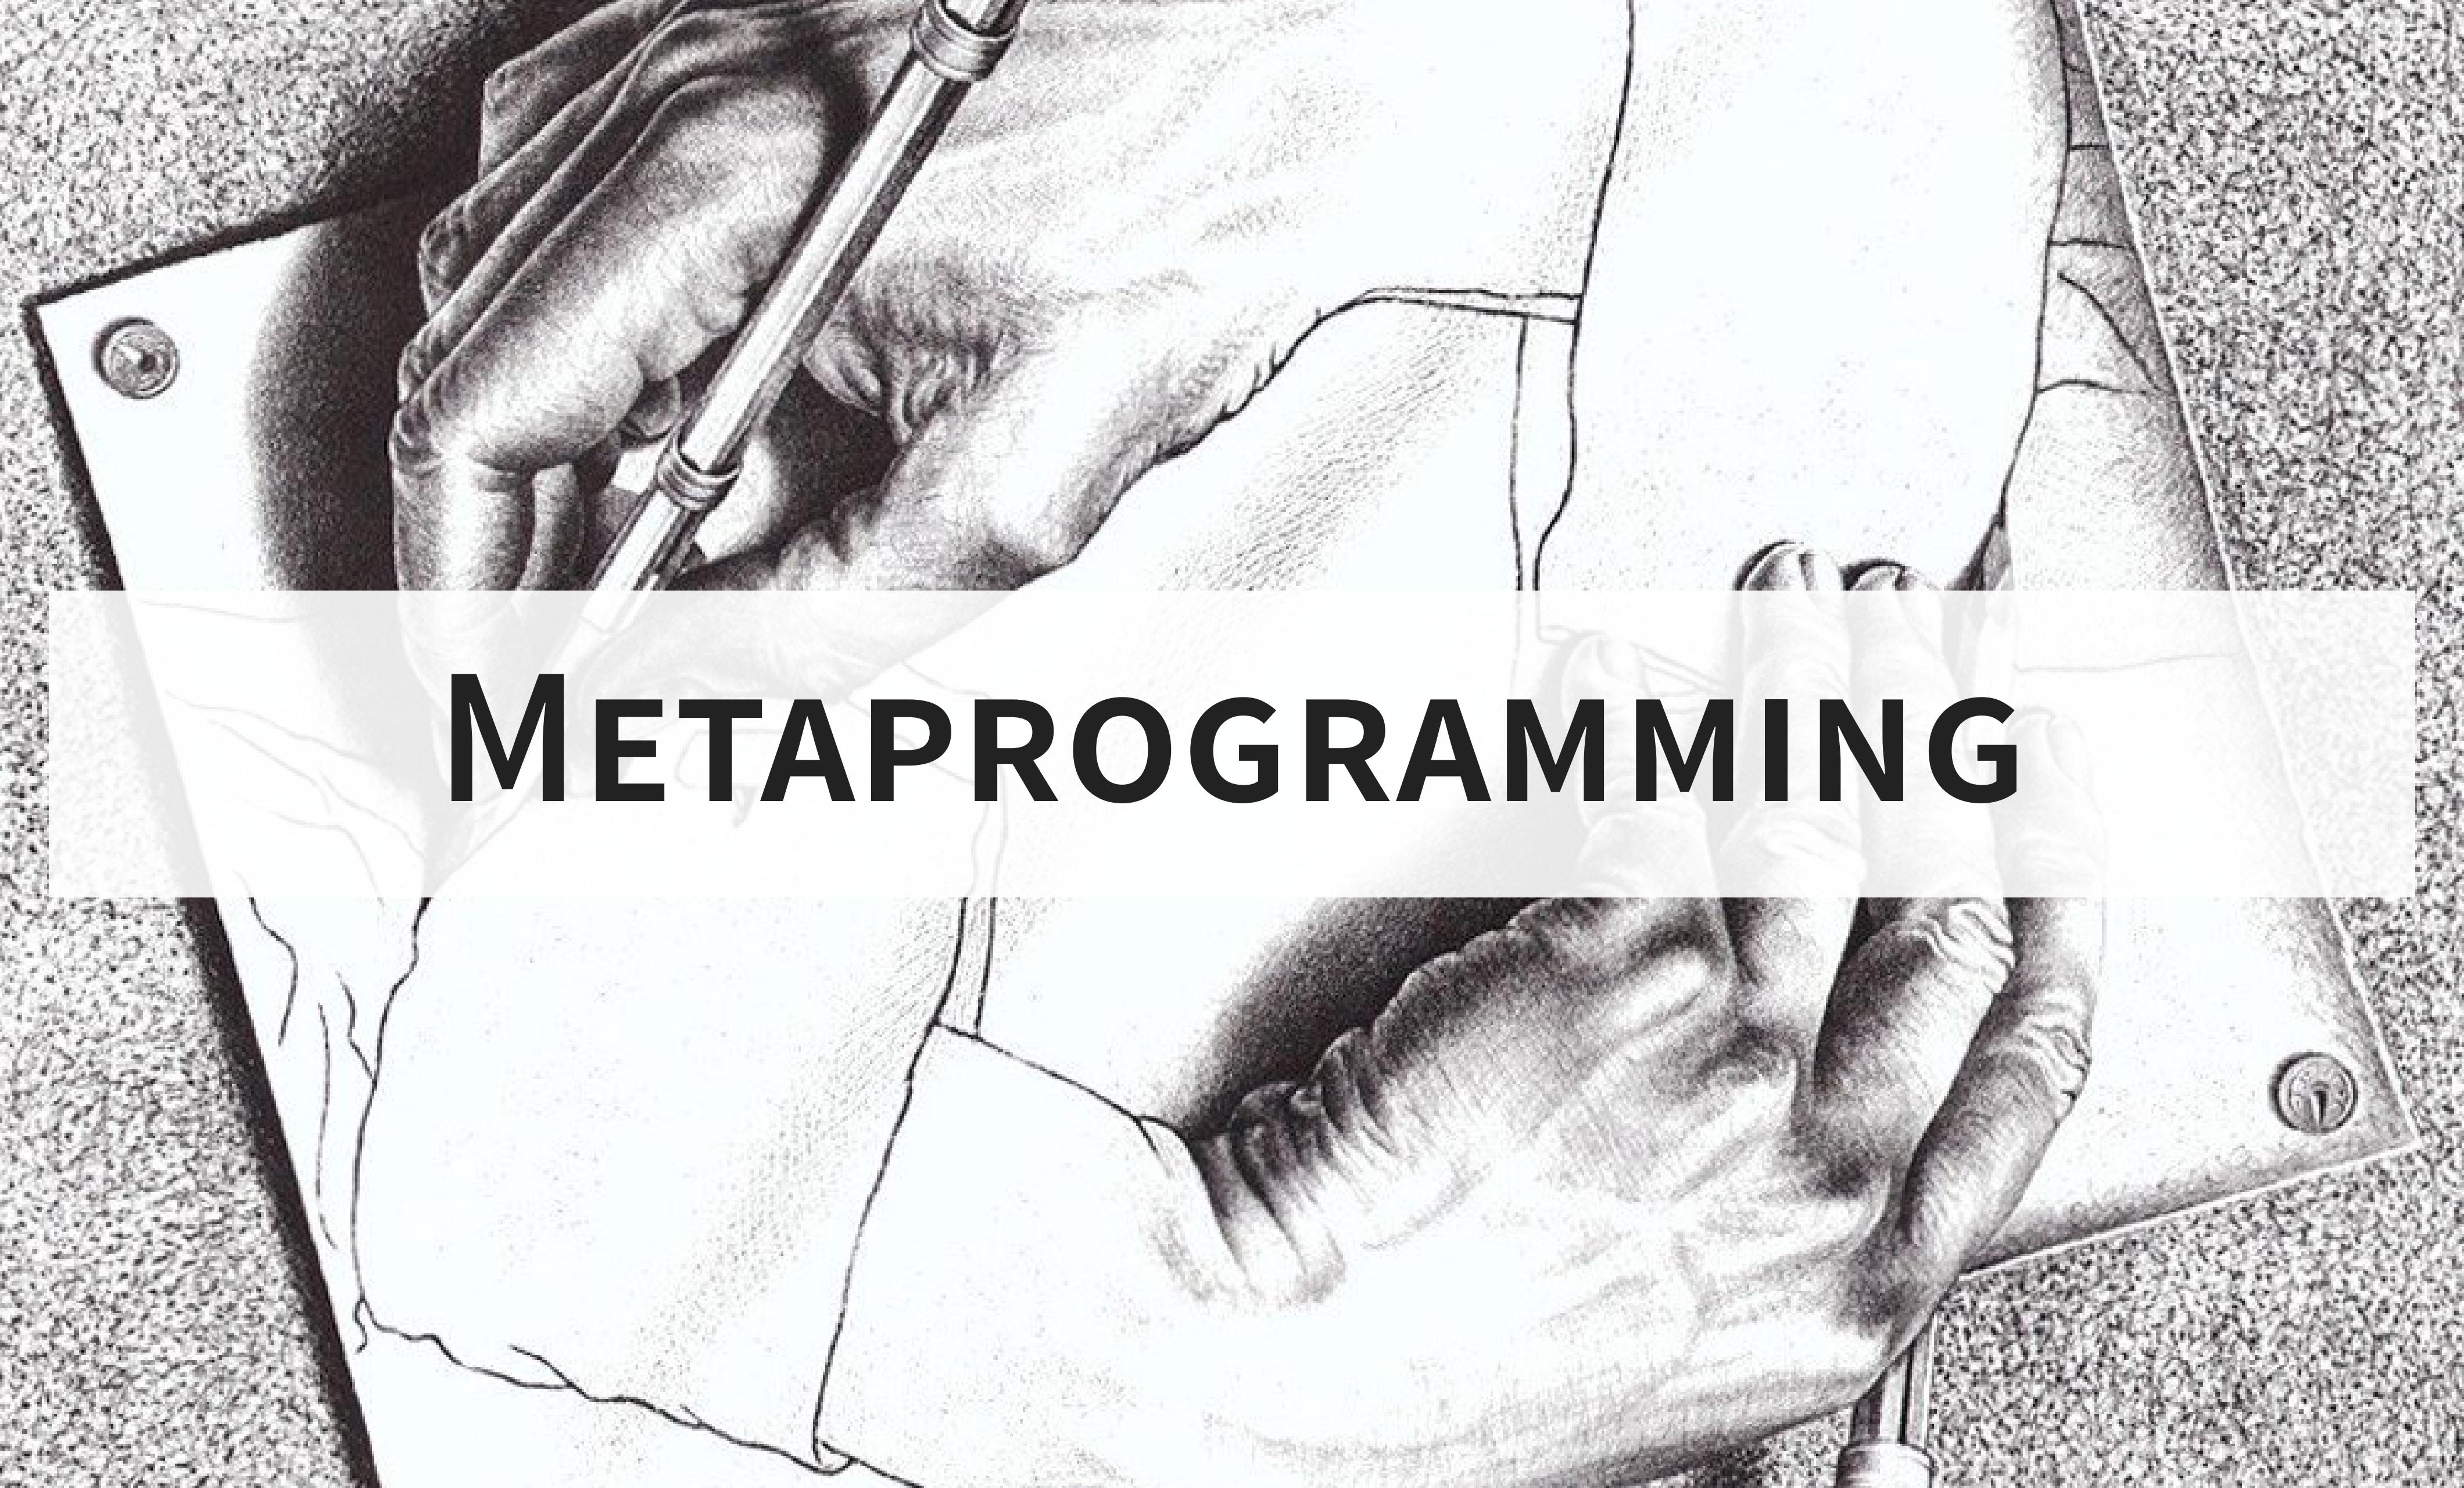
Dr. Scott Gorlin

Harvard University

Fall 2018

AGENDA

- Metaprogramming
- Atomic Targets
- Ser/Deser
- DB Design
- DB Juggling
- The Web
- API's
- Reading



METAPROGRAMMING

METAPROGRAMMING

Metaprogramming is a programming technique in which computer programs have the ability to treat programs as their data. It means that a program can be designed to read, generate, analyse or transform other programs, and even modify itself while running

— Wikipedia

WE'VE ALREADY SEEN THIS...

```
@unpack_argparse
def some_func(name='scott', **kwargs):
    # Old: some_func(args):
    #     print(args.name)
    ...

```

```
def unpack_argparse(func):
    @wraps
    def wrapped(args):
        return func(**{
            k: getattr(args, k)
            for k in dir(args)
            if not k.startswith('_')
        })
    return wrapped

```

```
if __name__ == '__main__':
    parser = ArgumentParser()
    parser.add_args(
        '-n', '--name',
        default='scotty')
    some_func(parser.parse_args())

```

A decorator is *so meta*.

WE'VE ALREADY SEEN THIS...

```
build([MyTask()])
dask_df.groupby().sum().compute()
```

When graph libraries modify the graph at runtime, they are meta programming

WE'VE ALREADY SEEN THIS...

```
# environment.yml
name: {{cookiecutter.project_slug}}
dependencies:
  - jupyter
{%- if cookiecutter.use_pytest == 'y' %}
  - pytest
{%- endif %}
  - python={{cookiecutter.python_version}}
  - setuptools_scm
```

Templated programs via jinja,
cookiecutter, etc

METACLASSES

Everything is an object, even classes!

```
class A:  
    def x(self):  
        return 5  
>>> a = A(); a.x()
```

... is the same as

```
A = type('A', tuple(), {'x': lambda self: 5})  
>>> a = A(); a.x()  
5
```

a is an instance of A, and A is an instance of type!

METACLASSES

We can subclass type to alter class creation!

```
class MyMeta(type):
    def __new__(cls, name, bases, dct):
        """Creates a new class!
        c = super().__new__(
            cls, name, bases, dct)
        c.x = lambda self: 5
        return c
```

```
class A(metaclass=MyMeta):
    pass

>>> a = A(); a.x()
5
```

This looks like it could be done with a decorator... except it runs *for all new subclasses too*

METACLASSES

```
class Timeable(type):  
  
    def __new__(cls, name, bases, dct):  
        dct.update({  
            'time_'+k: timeit(v)  
            for k, v in dct.items()  
            if callable(v)}  
    )  
    return super().__new__(  
        cls, name, bases, dct)
```

```
class MLModel(metaclass=Timeable):  
    def train(self, x, y):  
        ...  
  
>>> MLModel().time_train(x, y)
```

METACLASSES

```
class Property:  
    pass
```

```
class Encapsulated(type):  
    def __new__(cls, name, bases, dct):  
  
        props = [  
            k for k, v in dct.items()  
            if isinstance(v, Property)  
        ]  
        for p in props:  
            dct['get_' + p] = (  
                lambda self: getattr(  
                    self, '_' + p, None  
                )  
            )  
        ...
```

```
class Point(metaclass=Encapsulated):  
    x = Property()  
  
    >>> p = Point(); p.set_x(5)  
    >>> p.get_x()  
    5
```

METACLASSES

You likely don't need metaclasses (until you do)

But they are critical to understand complex code

MODEL INHERITANCE

```
class CommonInfo(Model):
    name = CharField()

    class Meta:
        abstract = True
```

```
class Student(CommonInfo):
    ...
```

... one Student table,
with all the columns

```
class Place(Model):
    name = CharField()

class Restaurant(Place):
    ...
```

... two tables;
Restaurant FK's to
Place because a
restaurant is a kind of
place

```
class Person(Model):
    name = CharField()

    def get_name(self):
        return self.name
```

```
class PigLatinPerson(Person):
    class Meta:
        proxy = True

    def get_name(self):
        n = self.name
        return ''.join([
            n[1].upper(), n[2:], 
            n[0].lower(), 'ay'
        ])
```

... just the Person
table

MODEL OPTIONS

```
class CommonInfo(Model):
    name = CharField()

    class Meta:
        abstract = True
```

```
class Student(CommonInfo):
    # Inherits `name`
    grade = FloatField()
```

How does django know Student
is a real table? It should have
inherited abstract = True!

MODELBASE

```
class ModelBase(type):
    def __new__(cls, name, bases, attrs):
        ...
        # Lots of magic for declarative options
        attr_meta = attrs.pop('Meta', None)
        ...

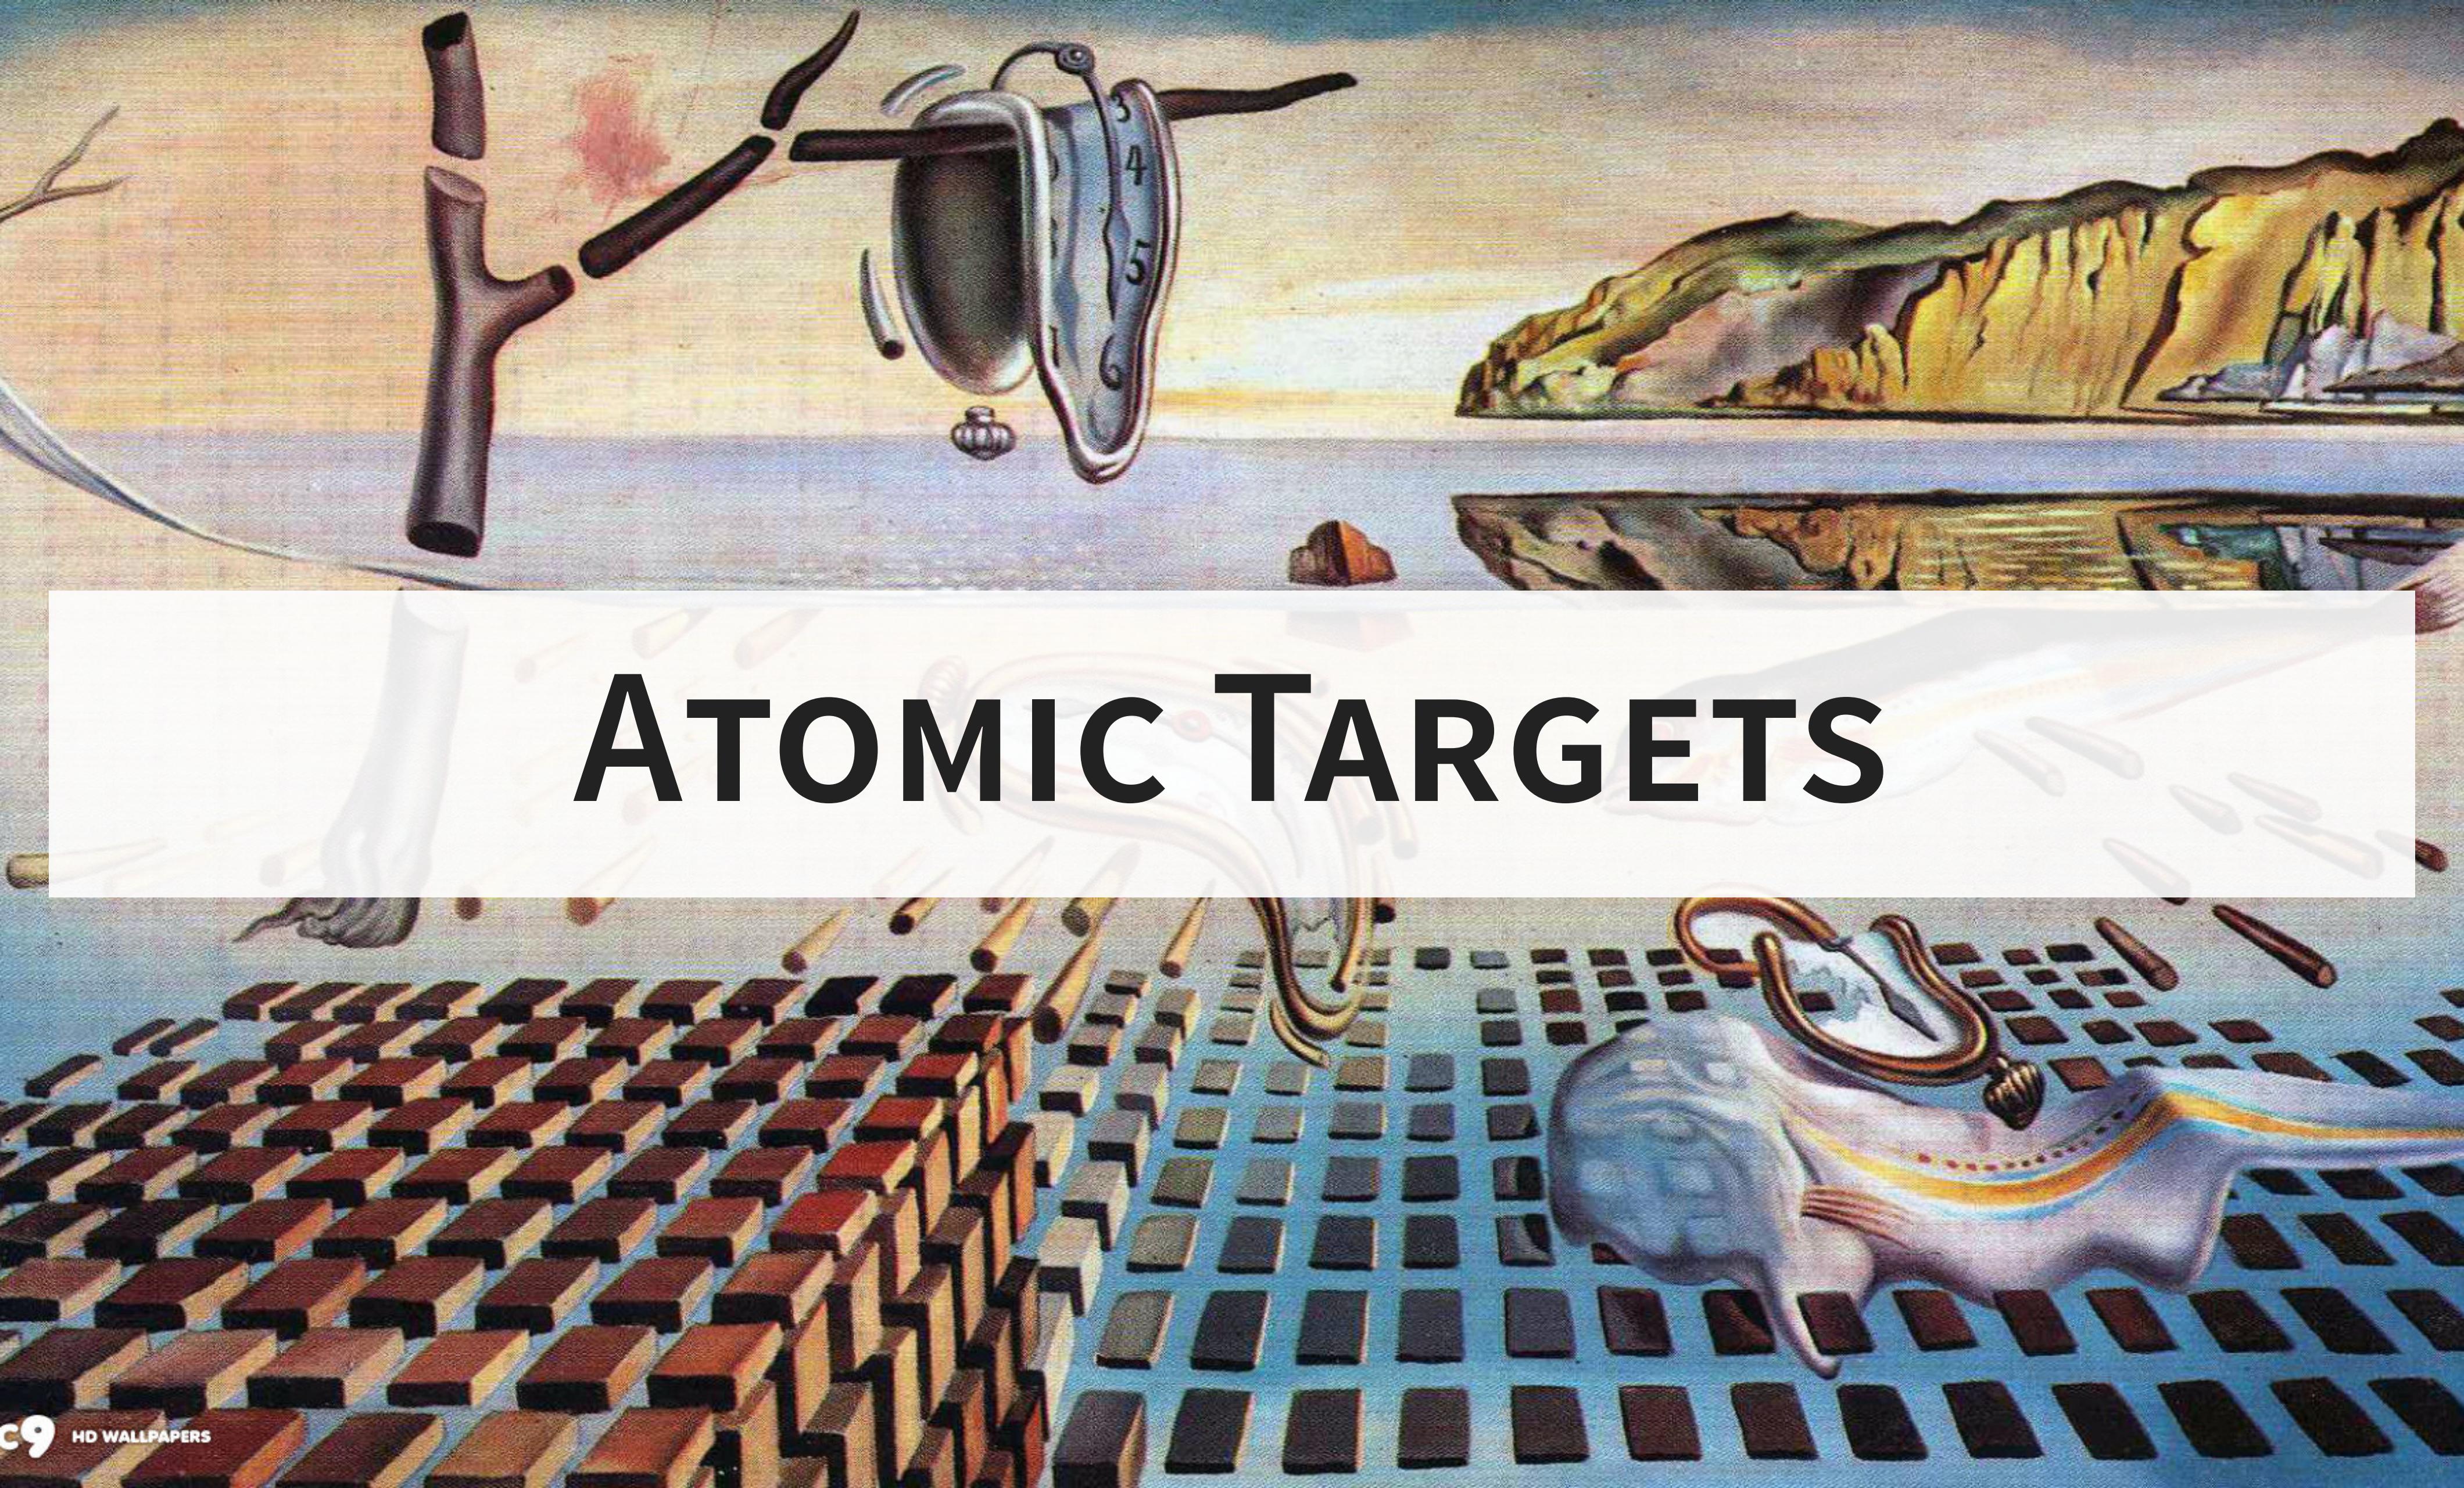
class Model(metaclass=ModelBase):
    ...
```

ModelBase extensively tweaks each new Model to ensure the fields work properly and to alter the standard inheritance patterns to be more intuitive for DB design

THE POINT

Metaprogramming allows you to write a Domain Specific Language, big or small

You can isolate the logic (the metaprogram) from the data (the actual program)



ATOMIC TARGETS

ATOMICITY

TRANSACTIONAL

```
with transaction.atomic():
    # NB: this is only for the default DB!
    MyModel.objects.create(...)
```

SQL gives us clean all-or-nothing transactions (usually)

WORKFLOW

```
class DjangoModelTarget(Target):
    def exists(self):
        # Check if row exists??
```

Luigi wants a target to complete atomically

DATABASE TARGETS

```
# luigi.contrib.postgres
class PostgresTarget(Target):
    def __init__(self, update_id, **conn_keywords):
        ...
        ...

    def exists(self):
        # Check *metadata*: did query run?
        "SELECT 1 from {marker_table}"
        " where update_id={self.update_id}"
```

Luigi has builtin sqlalchemy
Target's and Query tasks for
Postgres, Redshift, and generic
RDBMS

The update_id can be used to
implement a salted workflow!

Relying on an ‘existence’ table
violates a pure ‘can I read this
target’ form of atomicity

Can queries modify rows, or just
append, or must they support
both? What about unique keys?

DATABASE TARGETS

```
class DjangoModelTarget(Target):
    def __init__(self, model, **unique):
        ...
    def exists(self):
        try:
            # Assumes there is a unique
            # constraint
            m = self.model.objects.get(
                **self.unique_kwargs
            )
            return True
        except self.model.DoesNotExist:
            return False
```

You can easily write a target to check for existence of a unique (or unique_together) key

DATABASE TARGETS

```
class SaltedDjangoModelTarget(Target):
    def __init__(self, model, version, **unique):
        ...
    def exists(self):
        try:
            # Assumes there is a unique
            # constraint
            m = self.model.objects.get(
                version = self.version,
                **self.unique_kwargs
            )
            return True
        except self.model.DoesNotExist:
            return False
```

You can easily write a target to check for existence of a unique (or unique_together) key

You can even support salted targets

pk	name	version
1	Scott	8a3cef1
2	Scott	c091ef1

MULTIPLE Rows

```
class MultiTarget(DjangoTarget):
    def exists(self):
        return self.model.filter().count() > 0

    with atomic():
        target.model.bulk_create(...)
```

If any two targets have mutually exclusive rows, we can continue the pattern to multi-row, non-unique inserts

Except if a valid task should insert 0 rows, eg an empty dataset, and you want to know it is empty!

SER/DESER



PYTHON OBJECTS

Everything is a mutable dict!

```
class Point:  
    def __init__(self, x, y):  
        self.x = x  
        self.y = y  
  
p = Point(1, 2)  
p.adsf = 5 # Works just fine  
p.__dict__ # See inside
```

... unless it's a tuple or
implements __slots__

```
from collections import namedtuple  
  
Point = namedtuple('Point', ['x', 'y'])  
p = Point(1, 2)  
p.x # 1  
p.s = 14  
AttributeError: 'Point' object has no attribute 's'
```

PYTHON OBJECTS

Everything is a mutable dict!

Mutable types make for easy coding, at some overhead cost

Sometimes it is worth working with tuples for performance

```
>>> values = ['id', 'name']
>>> Blog.objects.values(*values)
<QuerySet [{'id': 1, 'name': 'Beatles Blog'}, ...]>
```

```
>>> Blog.objects.values_list(*values)
<QuerySet [(1, 'Beatles Blog'), ...]>
>>> DataFrame(_, columns=values)
      id          name
0    1  Beatles Blog
```

Note how the schema isn't replicated in every record!

DB DESIGN

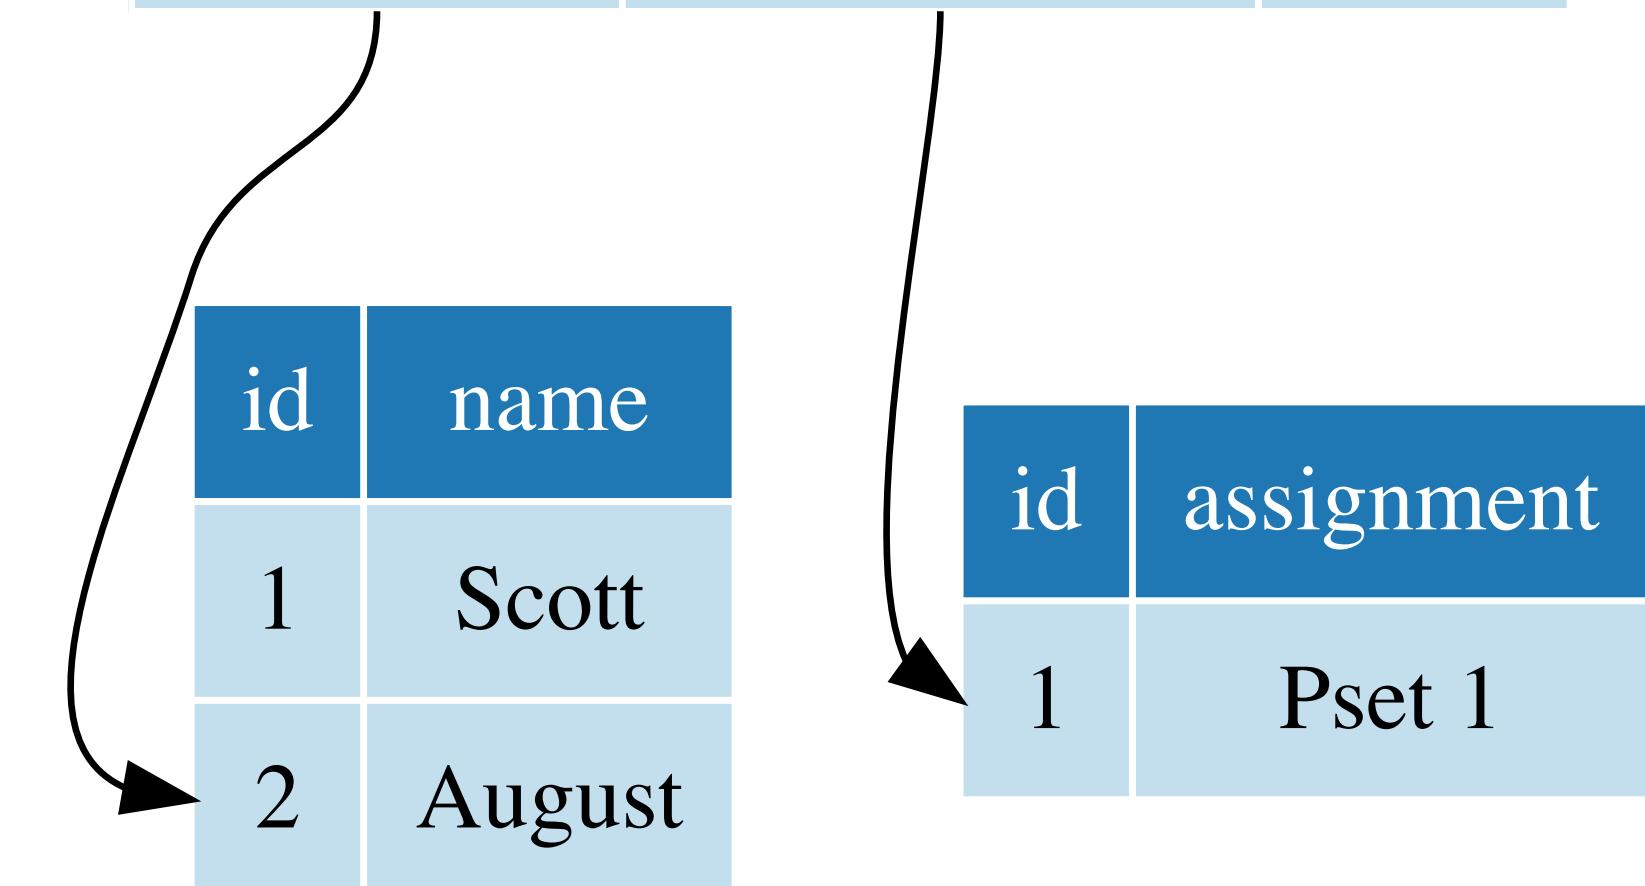
NORMALIZATION

DENORMALIZED

student	assignment	grade
Scott	Pset 1	89
August	Pset 1	91

NORMALIZED

student_id	assignment_id	grade
1	1	89
2	1	91



Normalized Forms reduce redundancy and may increase (or decrease) query complexity

NORMALIZATION

Normalization entails organizing the columns (attributes) and tables (relations) of a database to ensure that their dependencies are properly enforced by database integrity constraints.

— Wikipedia

NORMALIZATION

DENORMALIZED

```
class Grade(Model):  
    student = CharField()  
    assignment = CharField()  
    grade = FloatField()
```

```
Grade.objects.values(  
    'student'  
).annotate(grade=Avg('grade'))
```

Fast and easy, but how to
rename a pset?

NORMALIZED

```
class Assignment(Model):  
    assignment = CharField()
```

```
class Student(Model):  
    name = CharField()
```

```
class Grade(Model):  
    student = ForeignKey(Student)  
    assignment = ForeignKey(Assignment)  
    grade = FloatField()
```

How to query?

NORMALIZATION

NORMALIZED

```
class Assignment(Model):
    assignment = CharField()

class Student(Model):
    name = CharField()

class Grade(Model):
    student = ForeignKey(Student)
    assignment = ForeignKey(Assignment)
    grade = FloatField()
```

How to query?

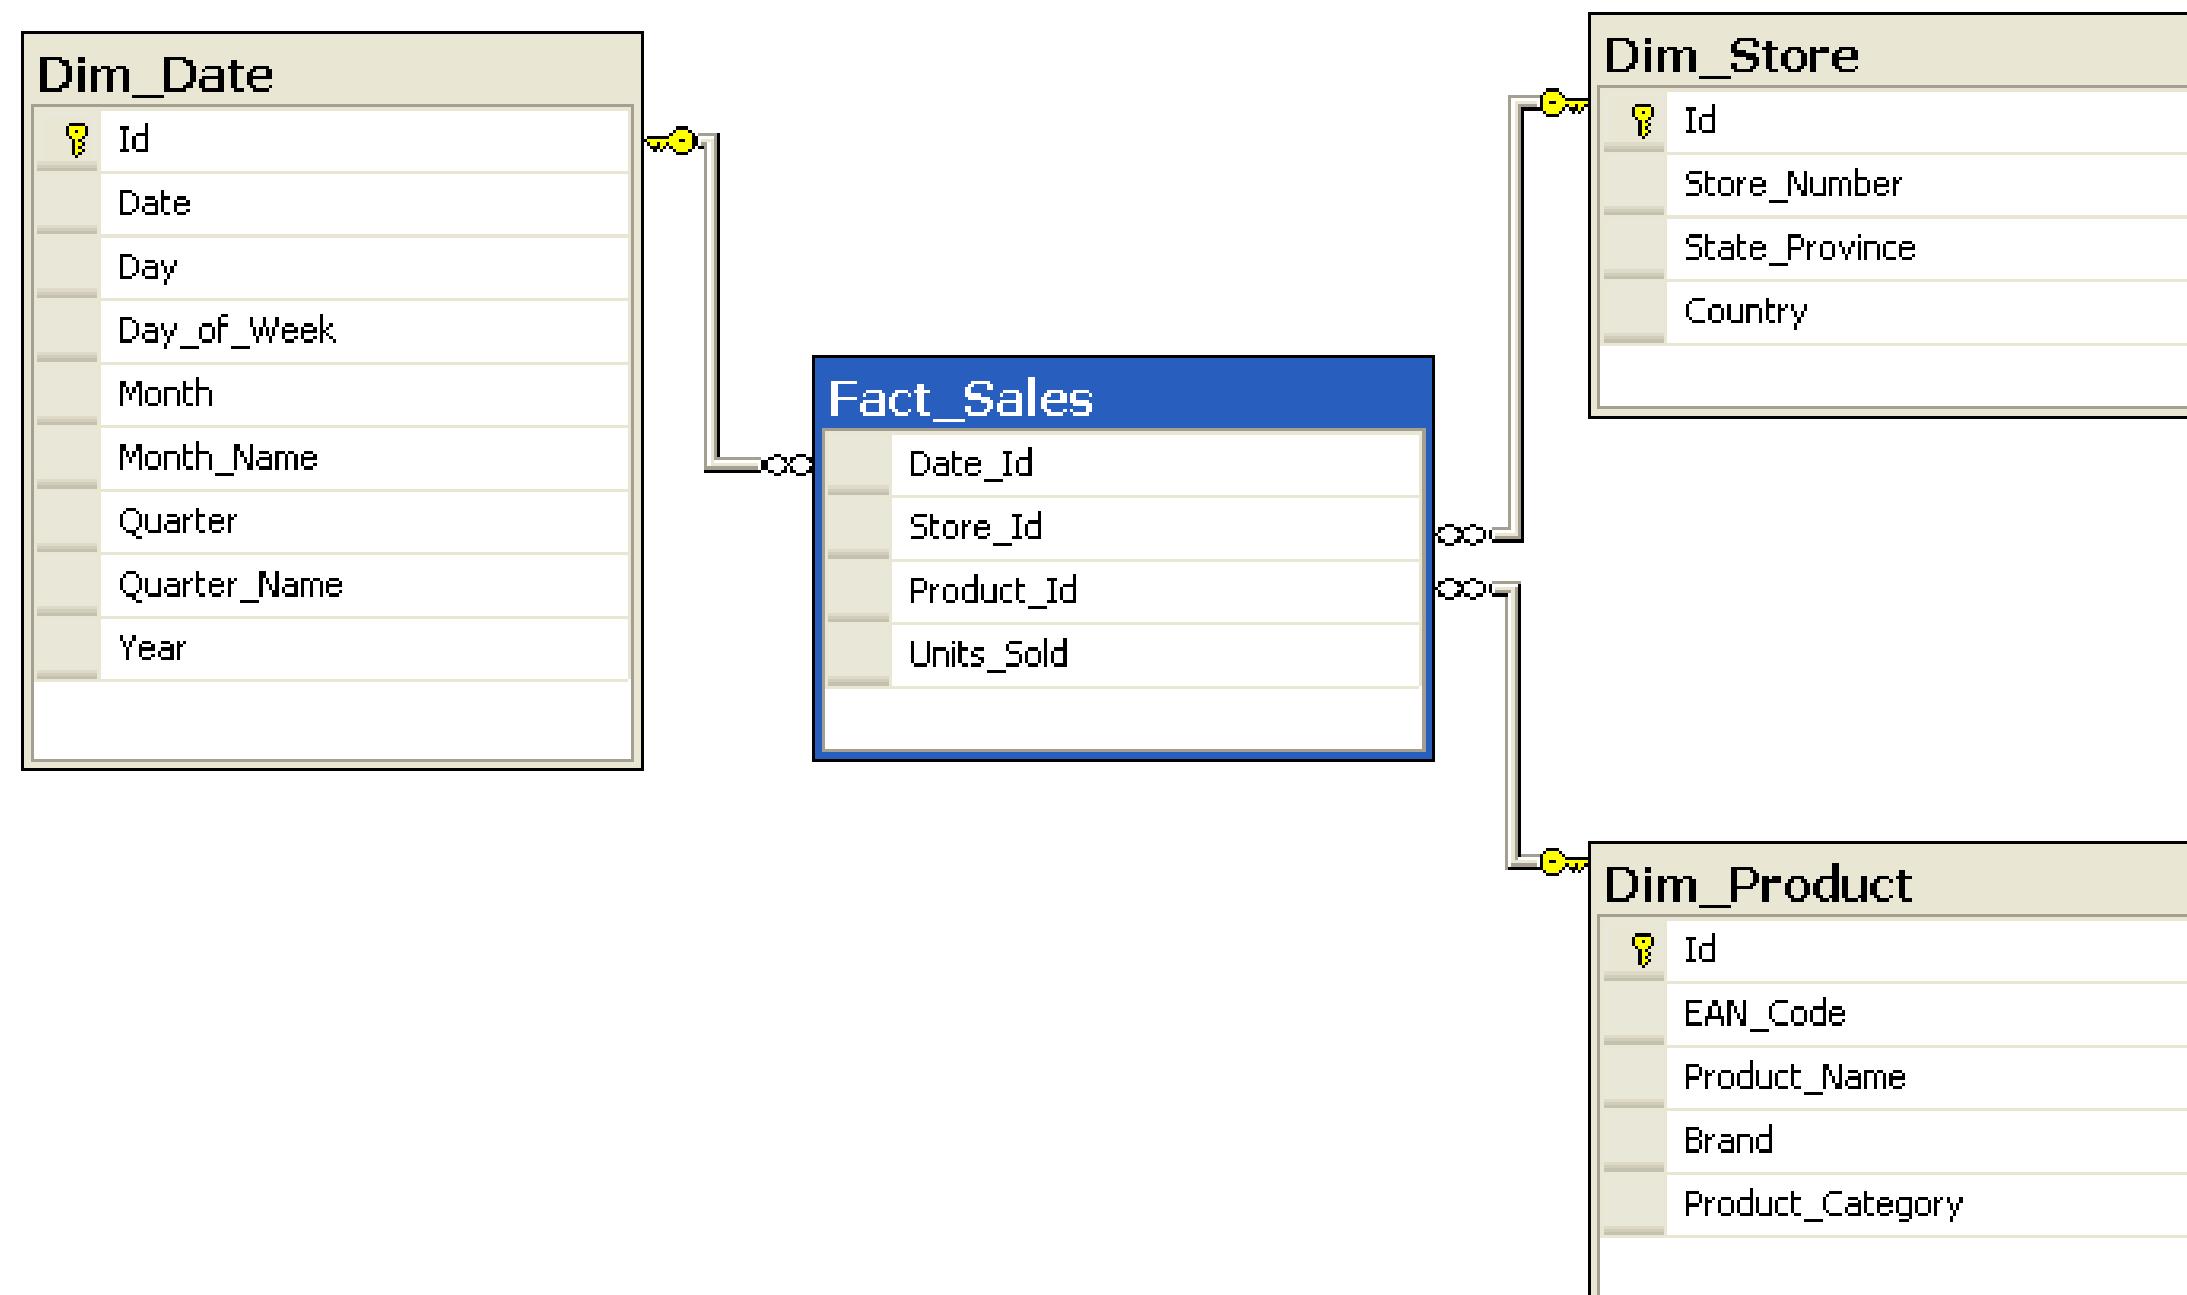
```
def v1(*args):
    return Grade.objects.values_list(*args)

# What you'd expect
>>> v1('student').annotate(grade=Avg('grade'))
[(Student('scott'), 89), ...]

# Fast and minimal, avoid join
>>> v1('student_id').annotate(...)
[(1, 89), ...]

# Join on 'real' value
>>> v1('student__name').annotate(...)
[('scott', 89), ...]
```

STARS AND FLAKES



STARS AND FLAKES

FACT TABLES

... are designed for aggregations

Every business ‘fact’ is captured in exactly one row (as an event, or as an aggregate)

Columns are summary stats (count, price, totals, etc)

DIMENSION TABLES

... are designed for groups and filters

Fact tables will FK to Dim tables to qualify rows

STARS AND FLAKES

FACT TABLES

```
class SalesFact(Model):
    # NB: Not DateField!
    date = ForeignKey(DateDim)
    product = ForeignKey(ProductDim)

    # Facts (theory should be just 1)
    units = IntField()
    revenue = FloatField()

total_sales = SalesFact.objects.sum()
```

DIM TABLES

```
class DateDim(Model):
    date = DateField()
    year = IntField()
    season = CharField()
    quarter = IntField()
    is_holiday = BooleanField()
```

Dims may store derived properties!

Compare:

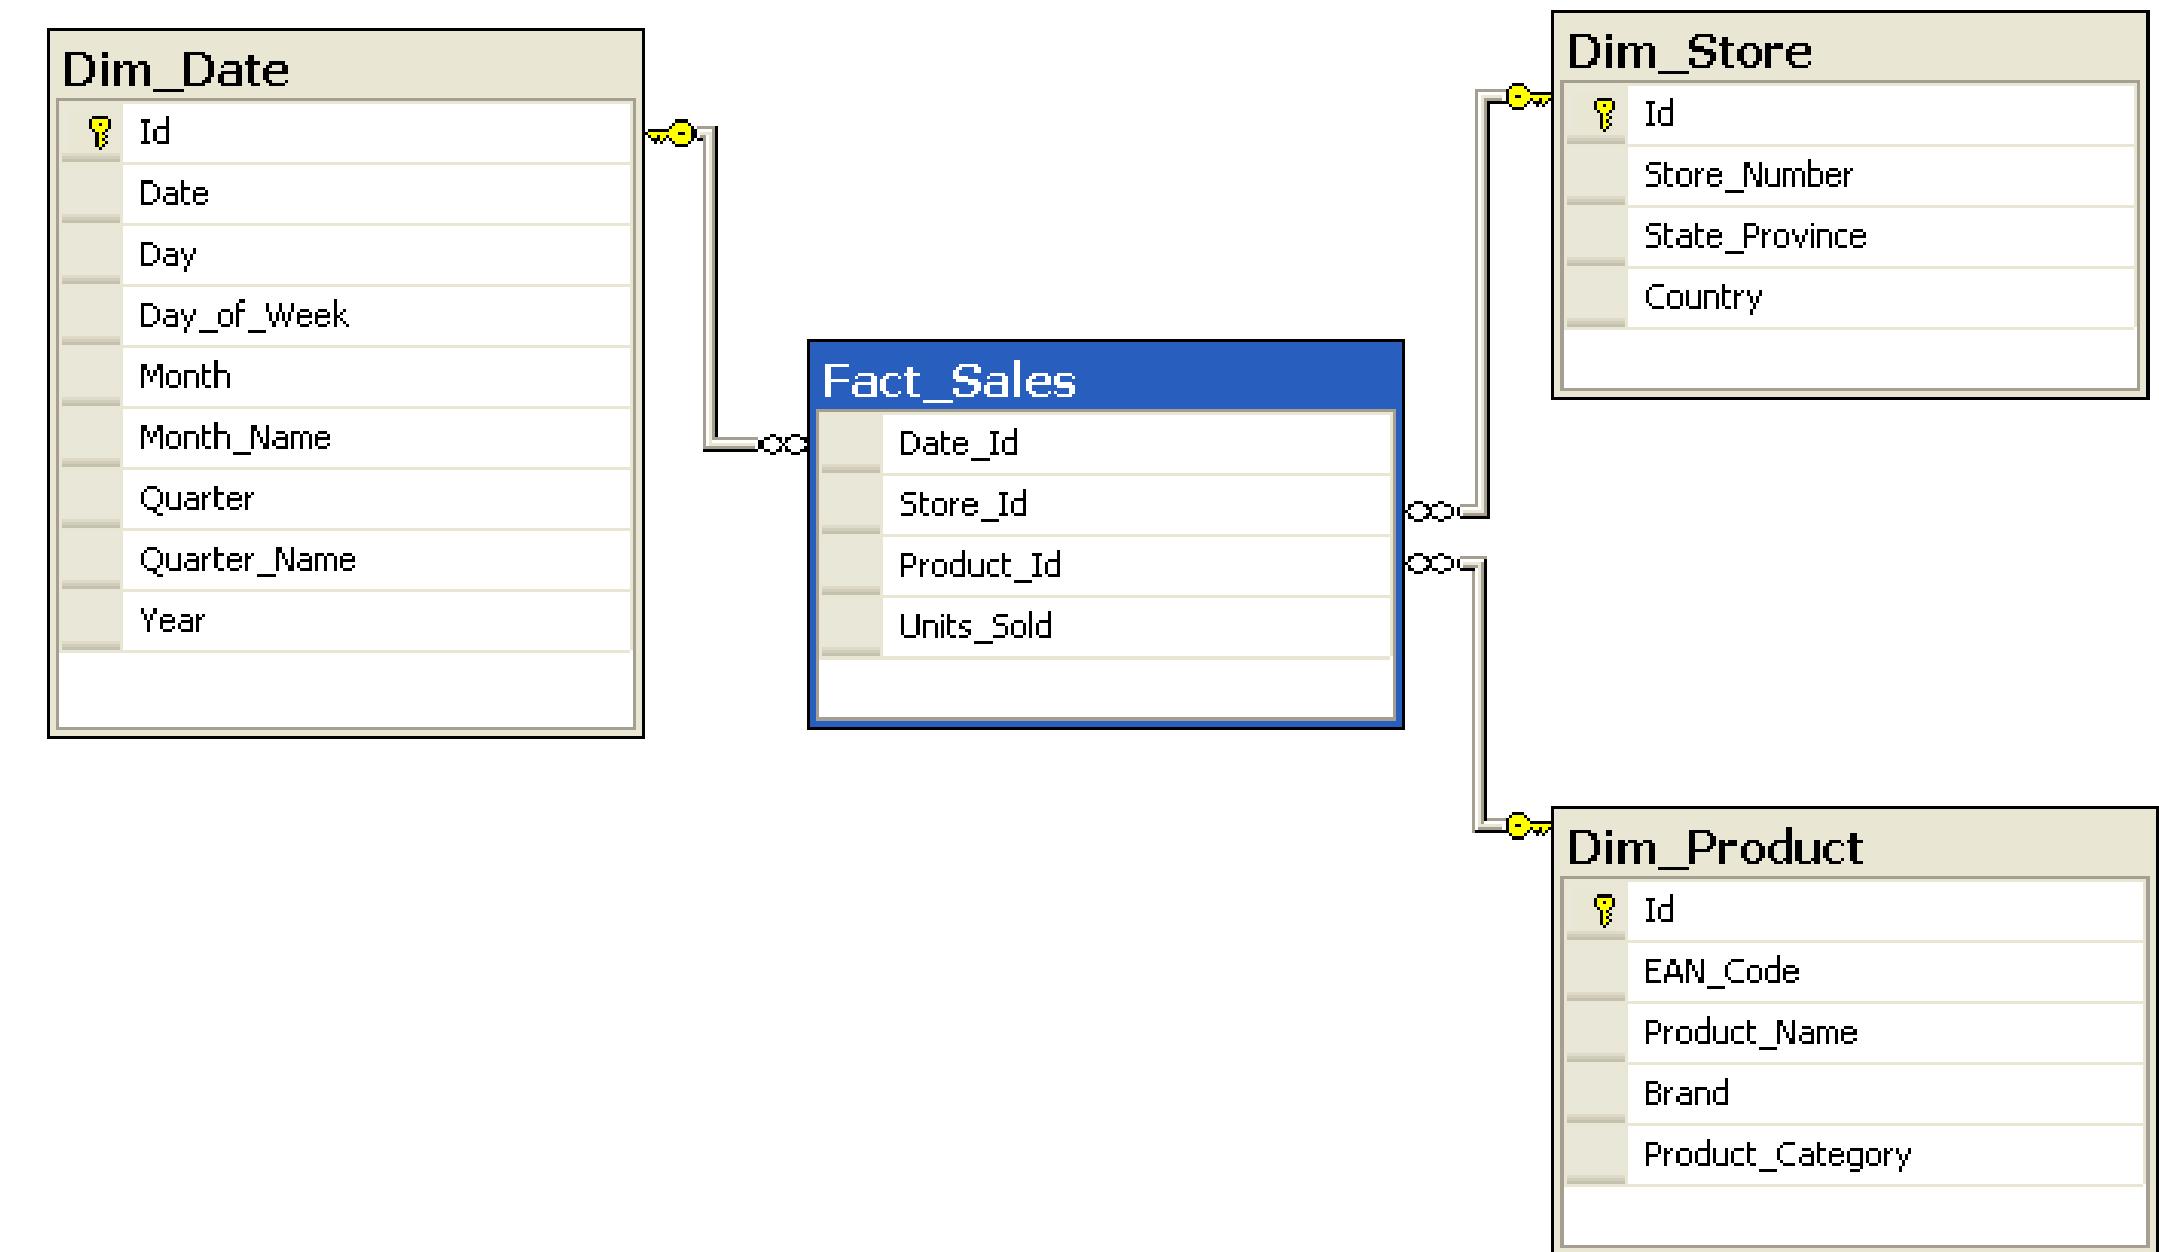
```
SalesFact.objects.filter(
    date__year=2018,           # Join and lookup, vs
    date__date__year=2018       # sql-side YEAR(date)
).sum()
```

STAR SCHEMA

Fact FK's to 1 level of Dim tables

Dim tables may have derivative properties to make easy querying

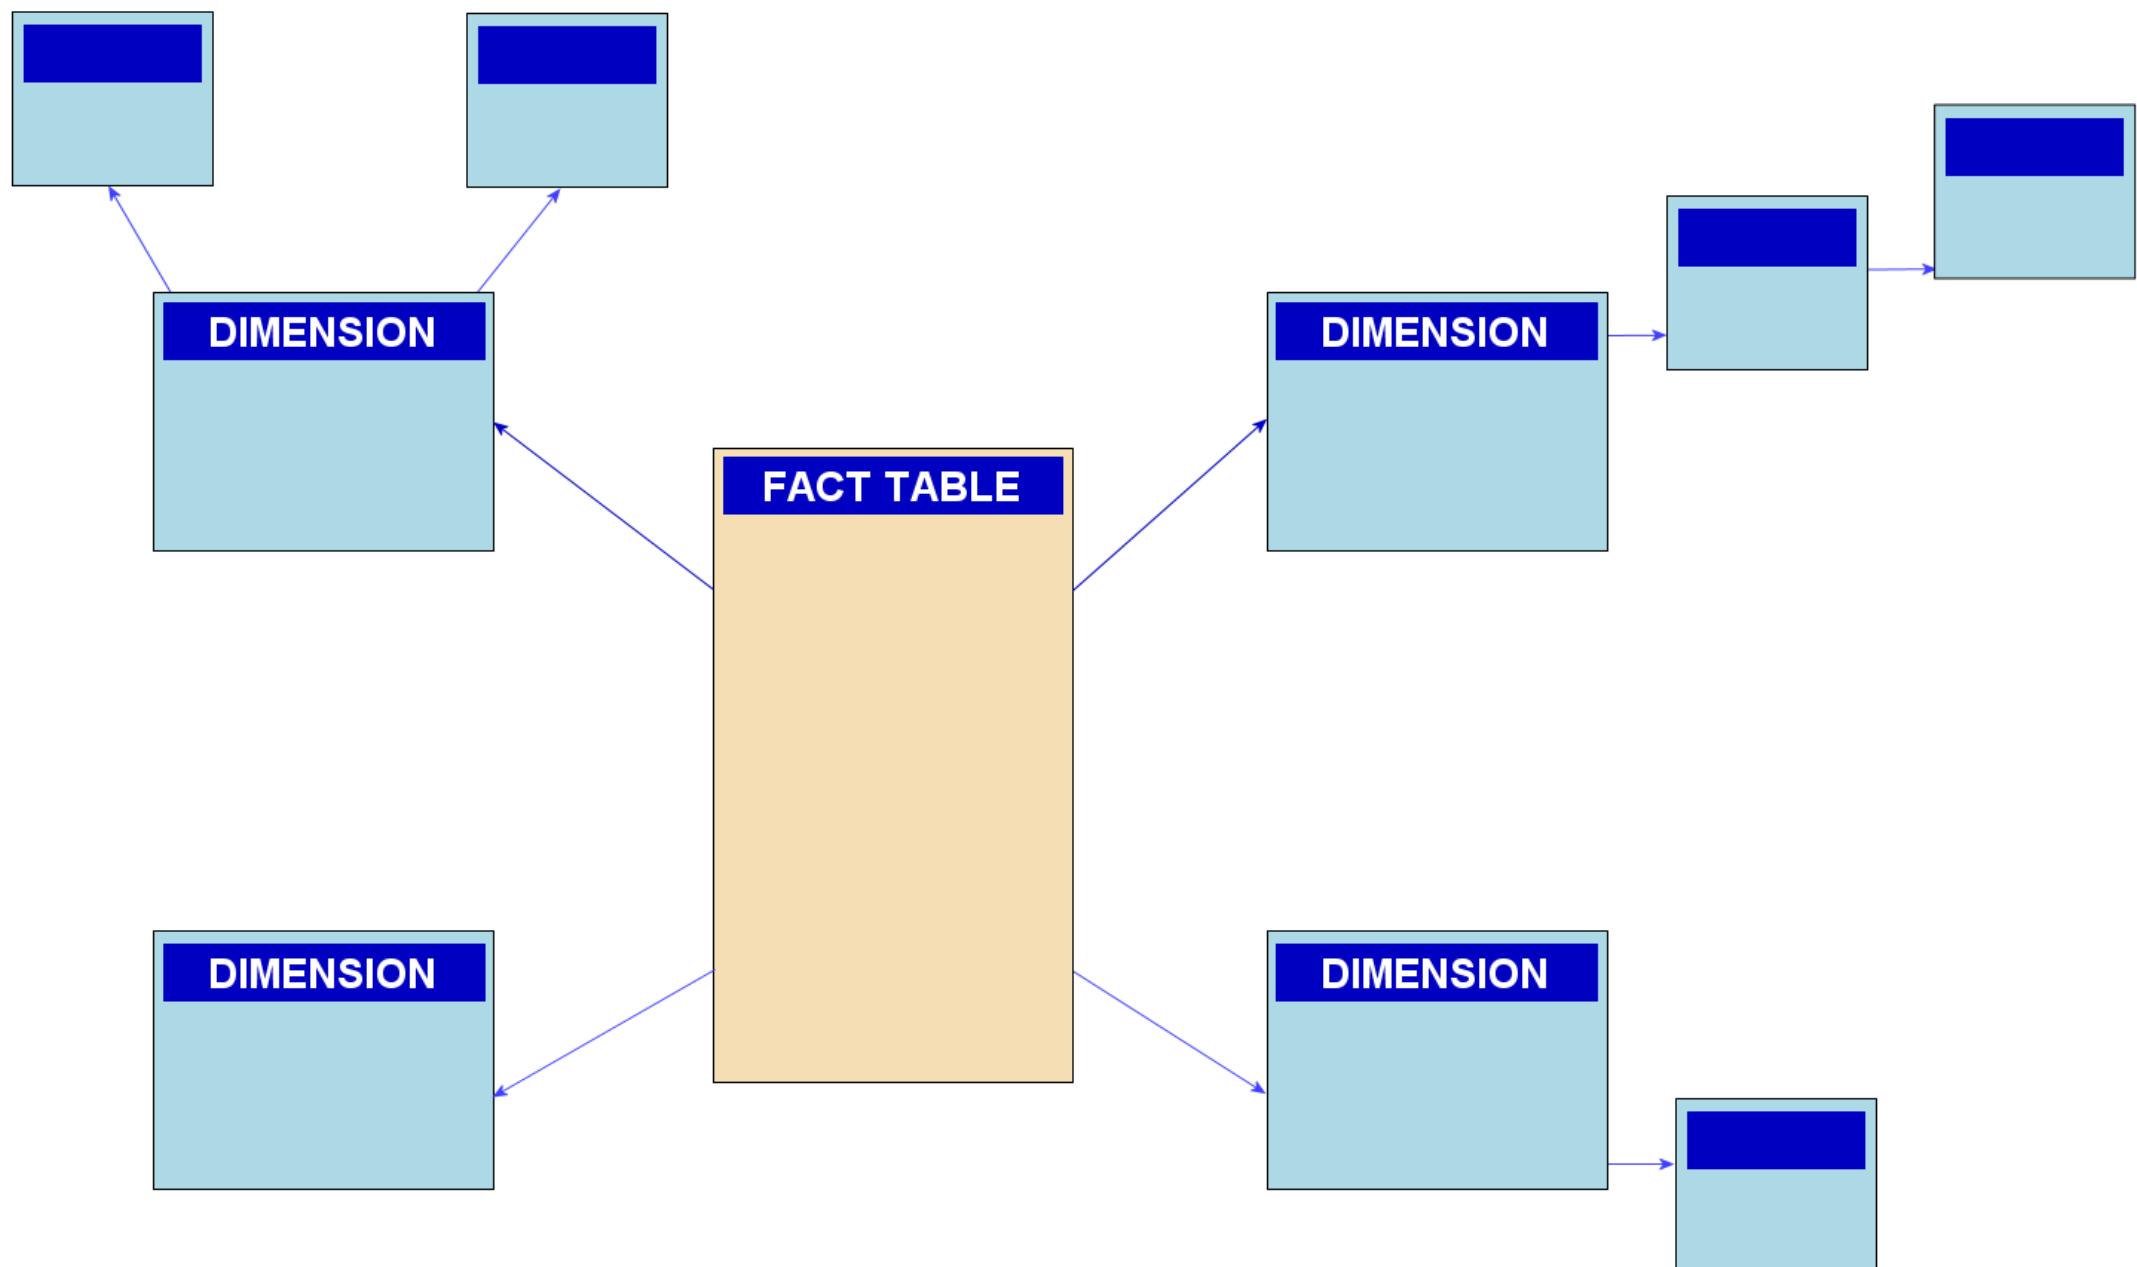
Common DataMart design due to fast writing



Star schemas are highly denormalized and usually don't guarantee unique dims! Avoid querying dim's by PK directly!

SNOWFLAKE SCHEMA

Fully normalized
dimensional variant of Star
More referential integrity,
but harder to query



INTELLIGENT KEYS

The primary key on a table is usually auto-incrementing and best considered ‘private’

But, Eng teams (and you) can optimize it

Beware over-optimization! This can kill readability

... or can be a critical performance boost, eg for *partitions* by date. SQL can't optimize queries across the join

```
class DateDim(Model):
    date = DateField()
    def save(**kwargs):
        if not self.pk:
            self.pk = self.get_pk(self.date)
        super(DateDim, self).save(**kwargs)

@classmethod
def get_pk(cls, date):
    return int(date.strftime('%Y%m%d'))
```

```
SalesFact.objects.filter(
    date_id=DateDim.get_pk(date.today())
).sum()
```

DB JUGGLING



MULTIPLE DB's

EXTERNAL DB's

It's common to rely on other DB's you don't own...

eg provided by Eng.

Try to keep *your* tables in one DB!

Write an isolated app for each DB

```
class ExternalTable(Model):
    class Meta:
        # You may set this to True
        # during unit tests!
        managed = False
```

MULTIPLE DB's

EXTERNAL DB's

Use *per-app routers* to configure!

Write an isolated app for each DB

```
class PerAppRouter(object):  
    """Specifies a DB for an entire app
```

In your `settings.py`, add::

```
ROUTING_TABLE = {  
    'my_app': 'my_db'  
}
```

"""

```
def db_for_read(self, model, **hints):  
    from django.conf import settings  
    return settings.ROUTING_TABLE.get(  
        model._meta.app_label, 'default')  
...
```

ATOMICS

Django assumes the default DB everywhere. If you're writing elsewhere, stick with sane defaults and patterns

```
def atomic_for_model(model, **kwargs):
    """Transactions for the db on a specific model

    Usage::

        with atomic_for_model(MyModel):
            ...

    """
    model_db = model.objects.all().db

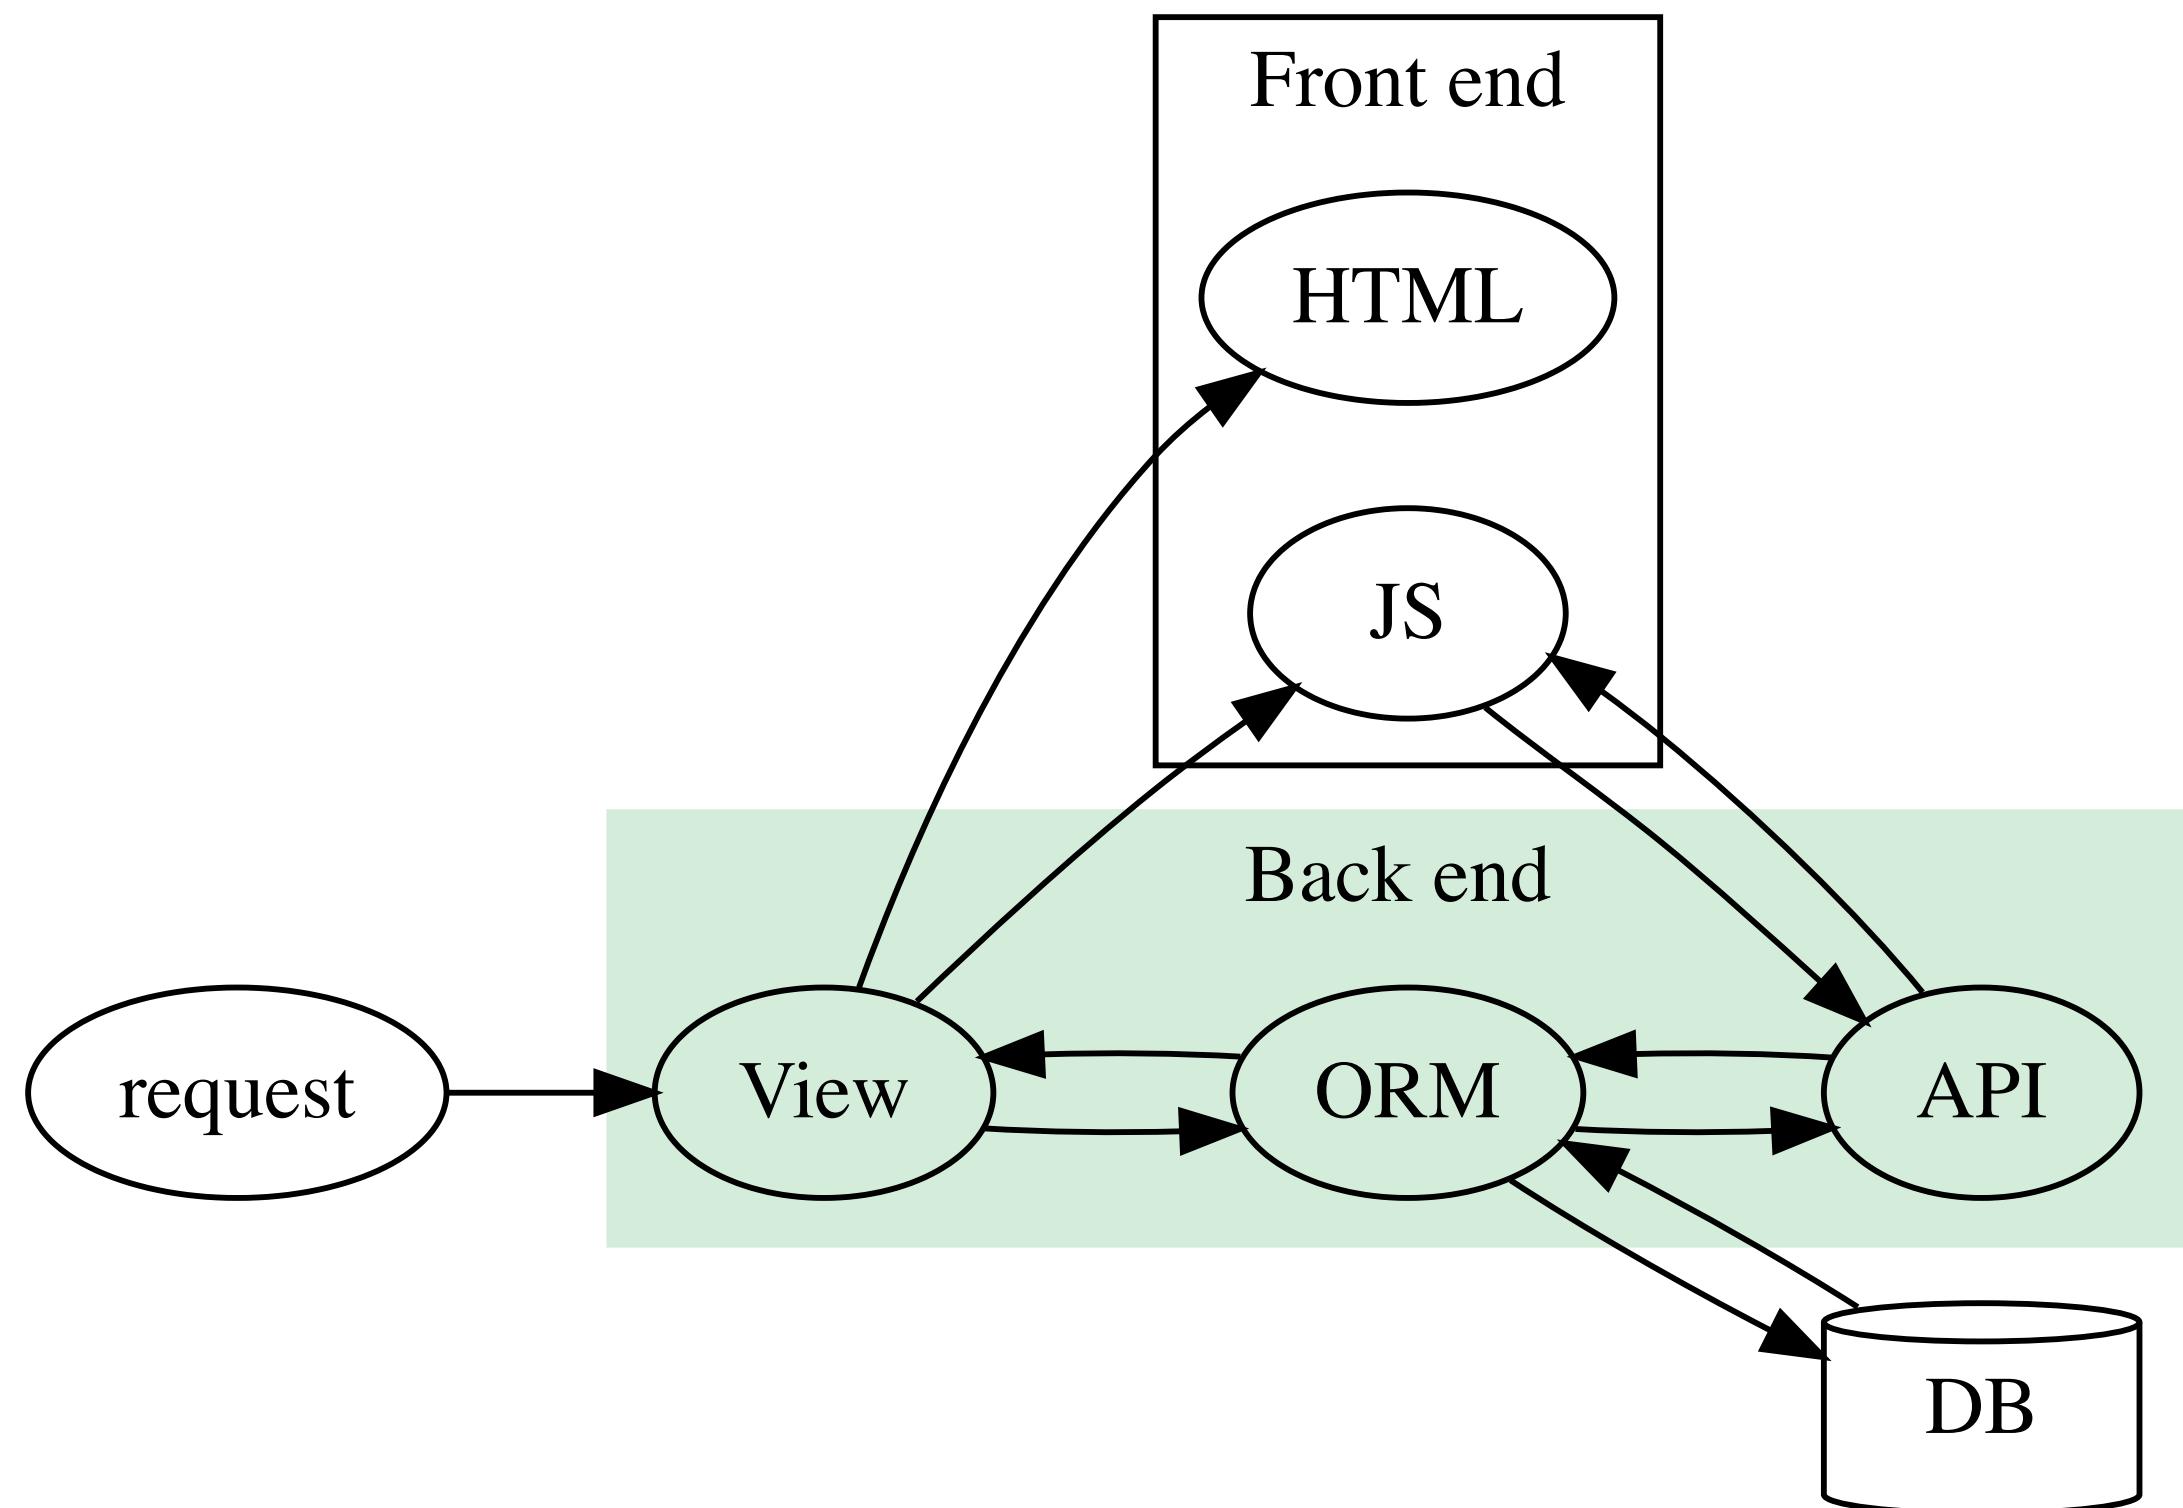
    return atomic(using=model_db, **kwargs)
```

http://WWW

THE WEB



PAGE LIFECYCLE



DJANGO VIEWS

URL ROUTING

```
from django.urls import path
from . import views

urlpatterns = [
    path('articles/<int:poll_id>', views.detail),
]
```

THE VIEW

```
def current_datetime(request):
    now = datetime.datetime.now()
    html = "<html><body>It is now %s.</body>
</html>" % now
    return HttpResponse(html)

def detail(request, poll_id):
    try:
        p = Poll.objects.get(pk=poll_id)
    except Poll.DoesNotExist:
        raise Http404("Poll does not exist")
    return render(request, 'polls/detail.html',
{'poll': p})
```

DJANGO VIEWS

It's easy to shoot out a quick web page w/ a template and some python data wrangling

Django provides a few other batteries like forms, some template tools, etc

Besides the built-in admin, Django does *not* provide a default frontend

BOOTSTRAP

Build responsive projects on the web with the world's most popular front-end component library.

Bootstrap is an open source toolkit for developing with HTML, CSS, and JS.

BOOTSTRAP

THE MISSING FRONTEND LIBRARY

The missing ‘make it look awesome’ frontend stack

<https://getbootstrap.com/>

I stole these alert themes from them!

Plus popups, layout, etc

BOOTSTRAP

Album example

Something short and leading about the collection below—its contents, the creator, etc. Make it short and sweet, but not too short so folks don't simply skip over it entirely.

Main call to action Secondary action

Thumbnail

This is a wider card with supporting text below as a natural lead-in to additional content. This content is a little bit longer.

View Edit 9 mins

Thumbnail

This is a wider card with supporting text below as a natural lead-in to additional content. This content is a little bit longer.

View Edit 9 mins

Thumbnail

This is a wider card with supporting text below as a natural lead-in to additional content. This content is a little bit longer.

View Edit 9 mins

Pricing

Quickly build an effective pricing table for your potential customers with this Bootstrap example. It's built with default Bootstrap components and utilities with little customization.

Free	Pro	Enterprise
\$0 / mo	\$15 / mo	\$29 / mo
10 users included 2 GB of storage Email support Help center access	20 users included 10 GB of storage Priority email support Help center access	30 users included 15 GB of storage Phone and email support Help center access
Sign up for free	Get started	Contact us

B © 2017

Features Resources About

Cool stuff Random feature Team feature

Resource Resource name Another resource

Team Locations Privacy

Company name Search

Dashboard

Orders Products Customers Reports Integrations

SAVED REPORTS Current month Last quarter Social engagement Year-end sale

20000 24000 22000 20000 18000 16000 Sunday Monday Tuesday Wednesday Thursday Friday Saturday

Section title

#	Header	Header	Header	Header	Header
1,001	1,001	1,001	1,001	1,001	1,001
1,002	1,002	1,002	1,002	1,002	1,002
1,003	1,003	1,003	1,003	1,003	1,003
1,004	1,004	1,004	1,004	1,004	1,004

... BUT!

OLD SCHOOL COOL

Server-generated webpages are so 2005

... even if written in python!

You don't want Chrome to sit there spinning while the ORM works!

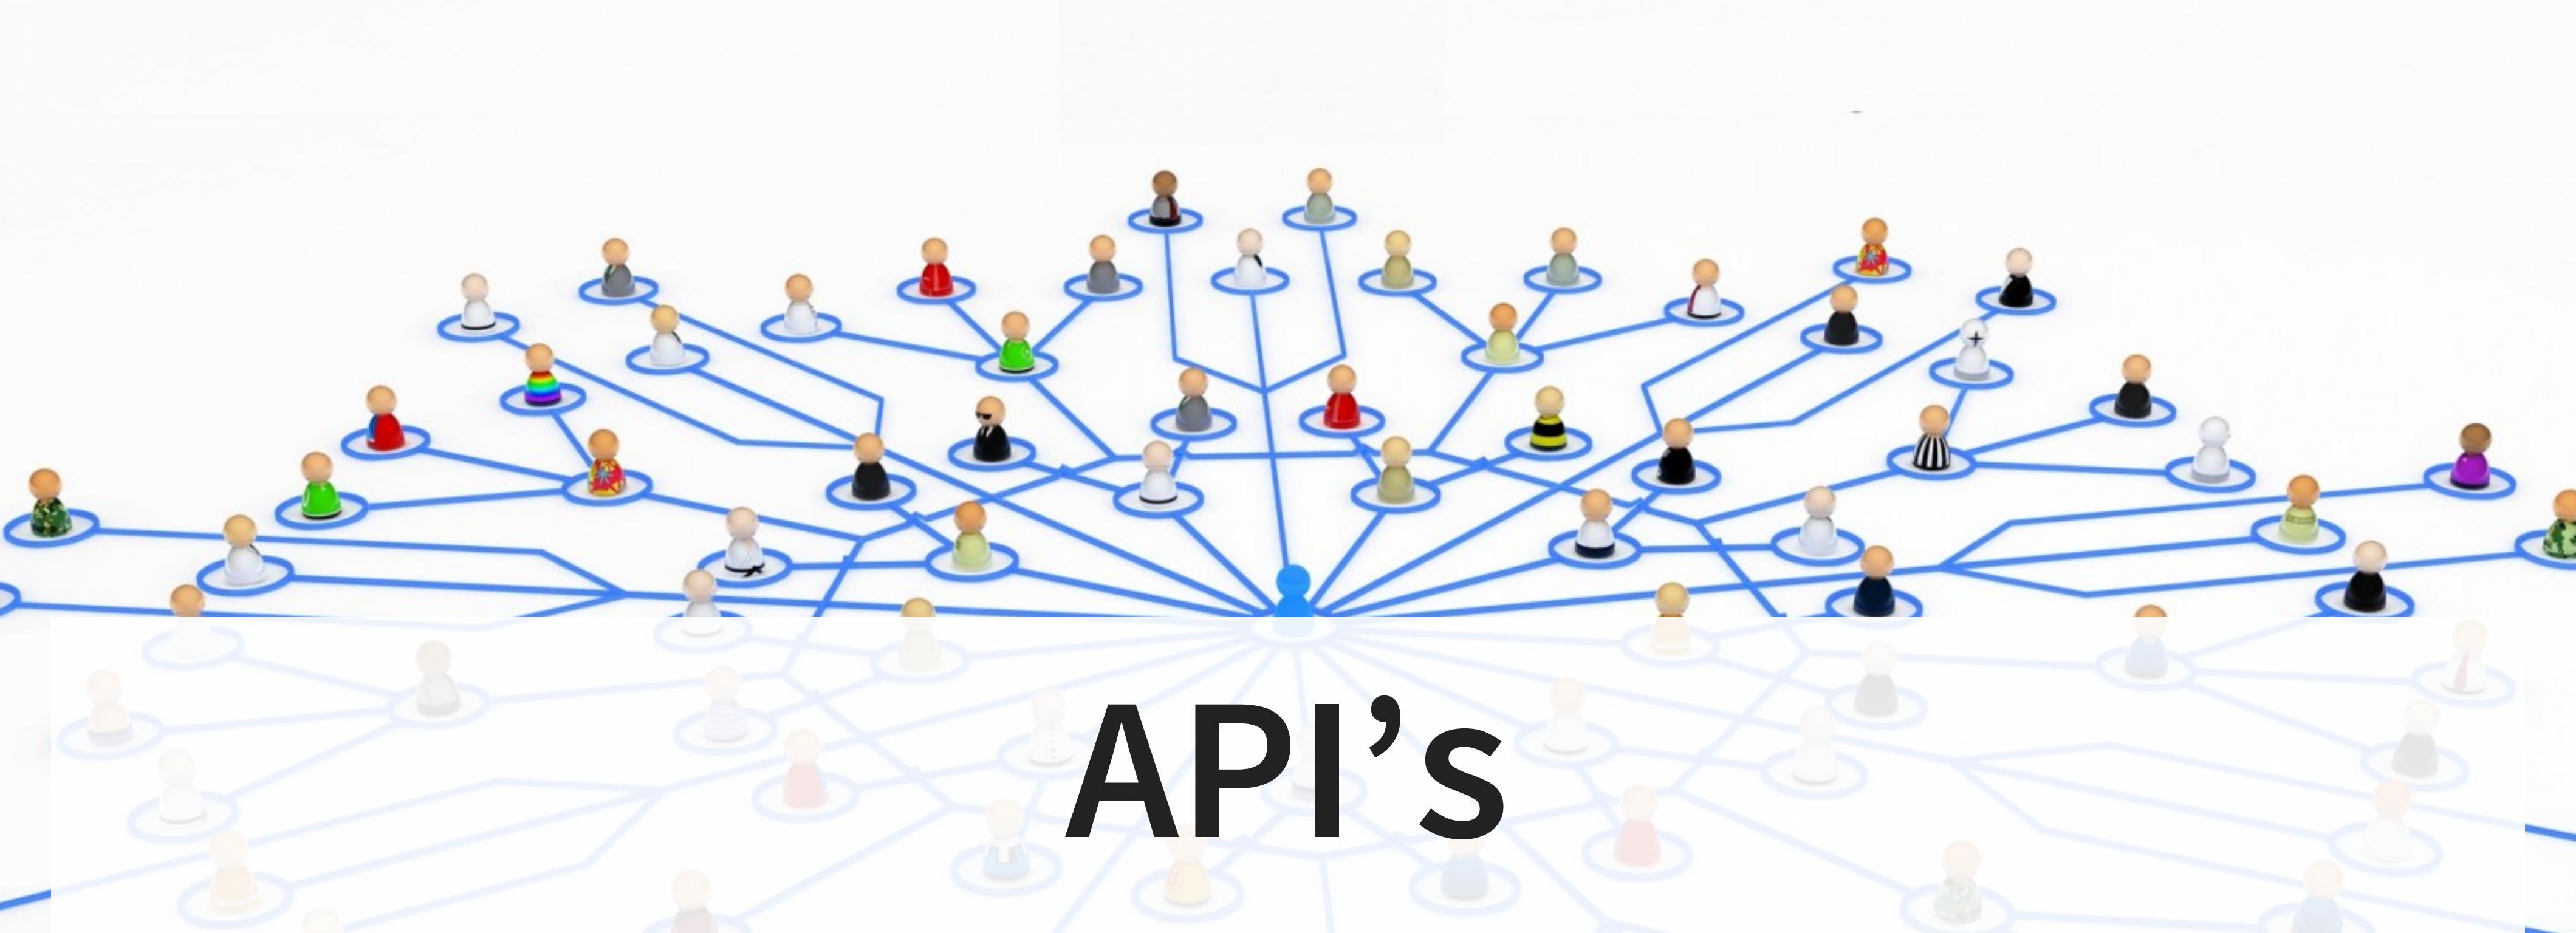
... though they do work great for some things

RESPONSIVE SITES

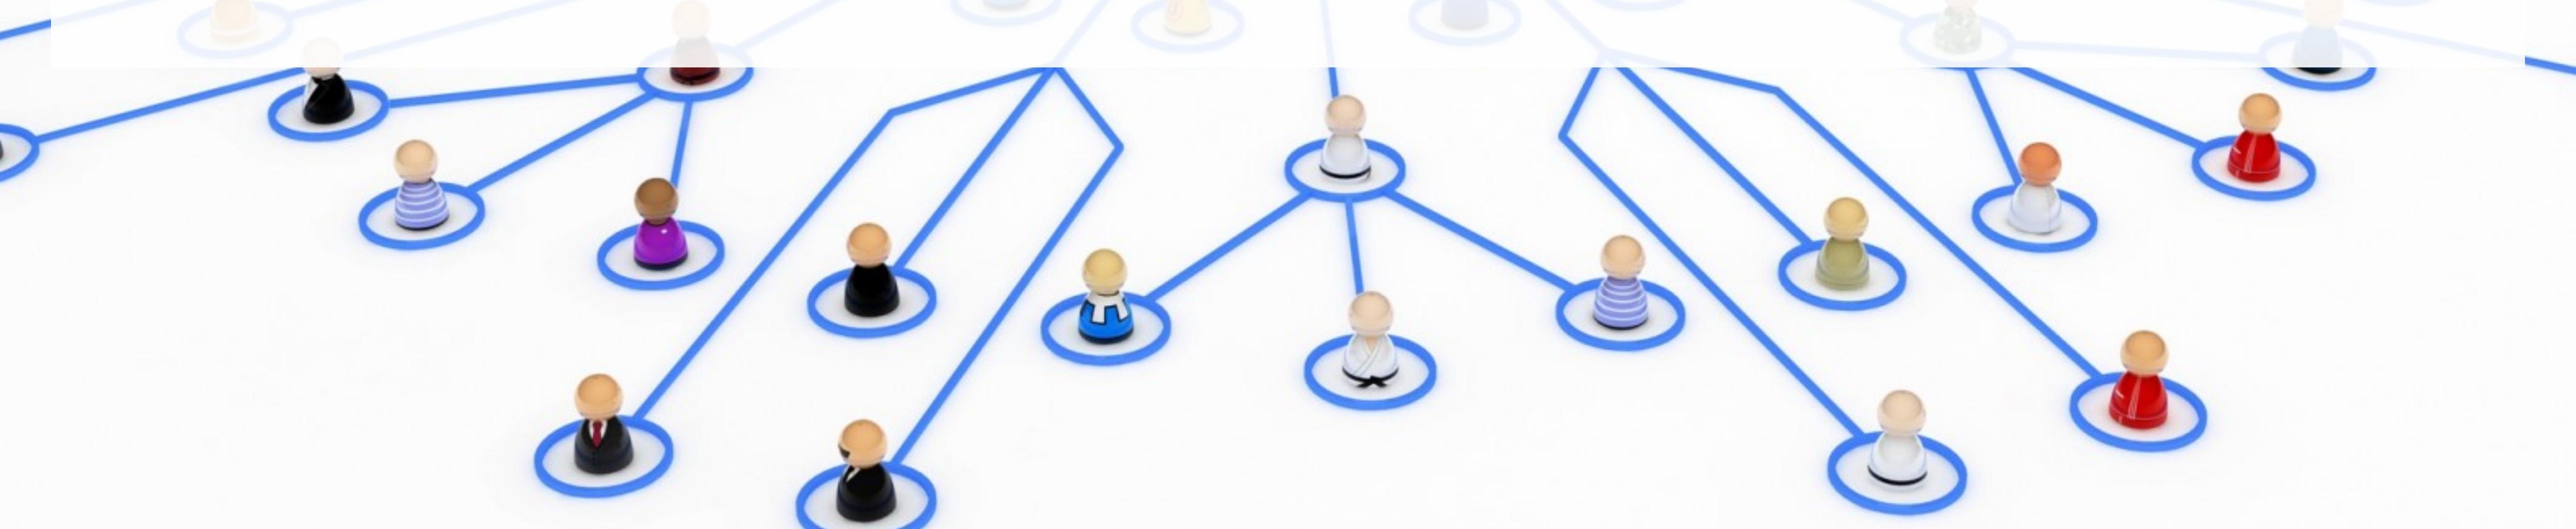
Even as a DS, you want a web view that is live and reactive

It may hit multiple DB's, load big slow data...

You may want interactivity! You need *javascript and API's!*



API'S



ORM-AS-A-SERVICE

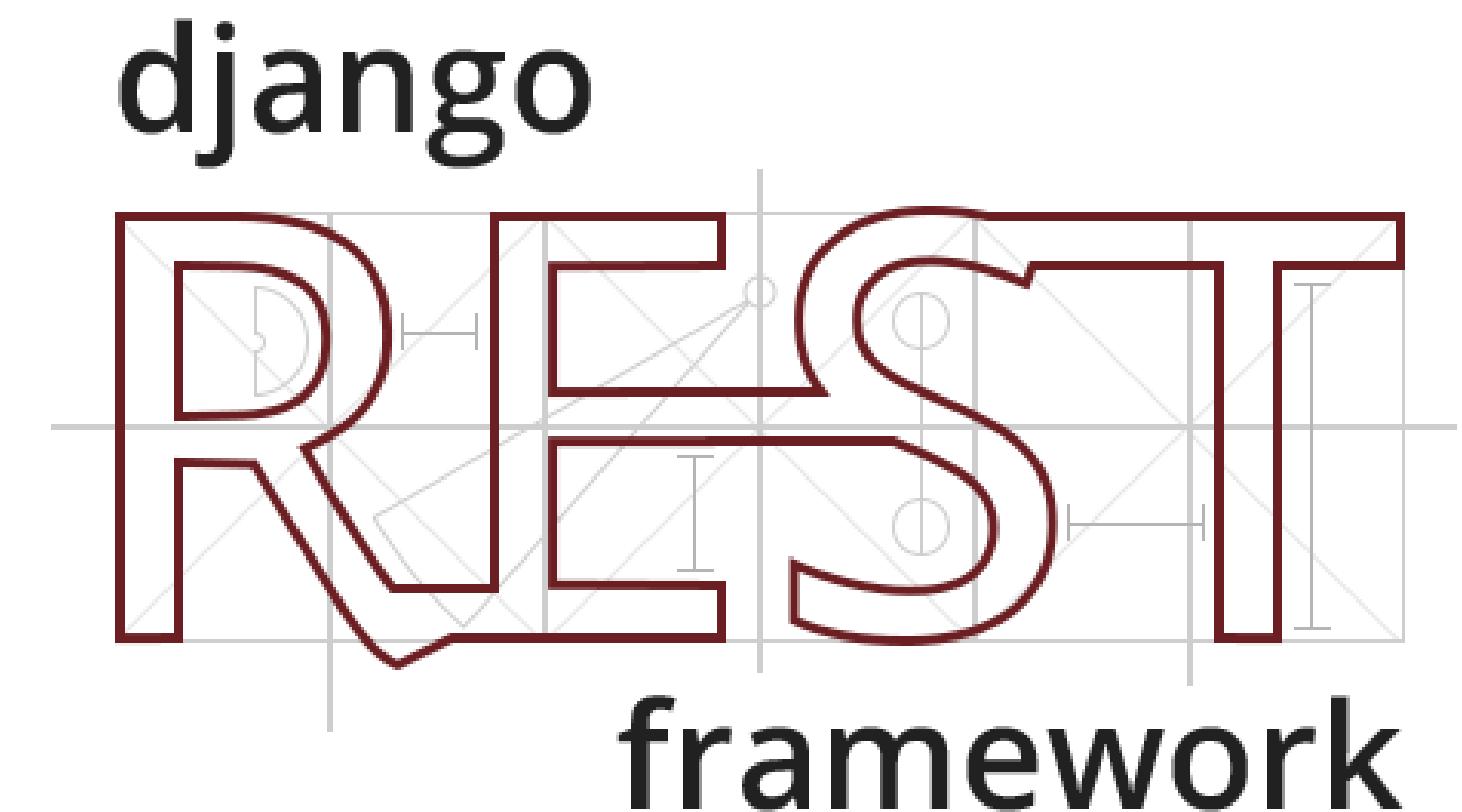
A django view hard-bakes the data into the page

You want to expose your data directly to clients, whether for the web via JS, or any other service (even python)

DJANGO REST FRAMEWORK

Your first stop to expose your
data...

<https://www.djangoproject.com/>

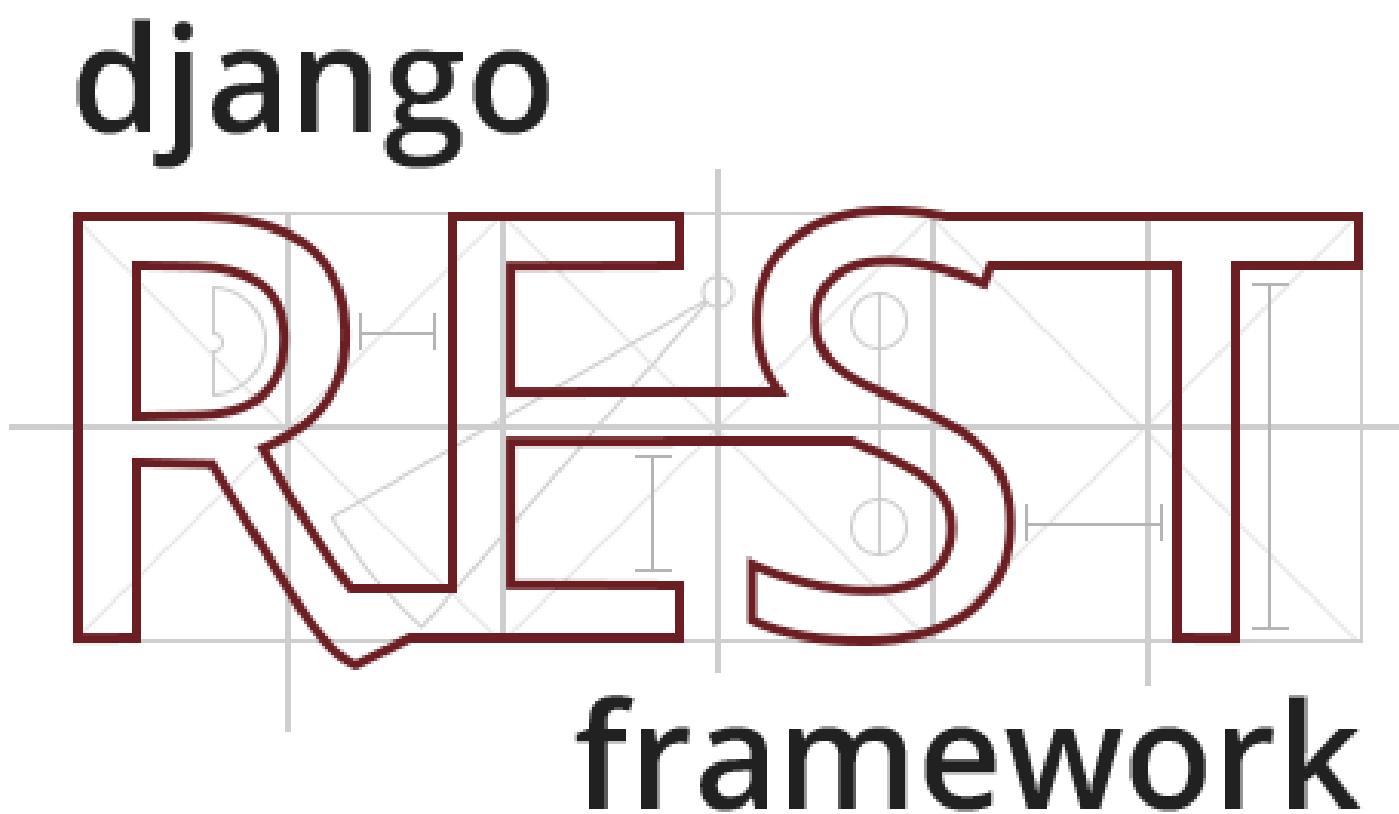


DJANGO REST FRAMEWORK

It's not really REST, but that's okay

It emphasizes multiple linked endpoints for 'raw' data access

It emphasizes development and debugability



SERIALIZERS

It's easy to expose any Django model...

```
class AccountSerializer(serializers.ModelSerializer):
    class Meta:
        model = Account
        fields = '__all__'
```

SERIALIZERS

... or any other object

```
class CommentSerializer(serializers.Serializer):
    email = serializers.EmailField()
    content = serializers.CharField(max_length=200)
    created = serializers.DateTimeField()
```

VIEWS

... and pack it all into a view

```
class SnippetList(APIView):

    def get(self, request, format=None):
        snippets = Snippet.objects.all()
        serializer = SnippetSerializer(
            snippets, many=True)
        return Response(serializer.data)

    def post(self, request, format=None):
        serializer = SnippetSerializer(
            data=request.data)
        if serializer.is_valid():
            serializer.save()
            return Response(
                serializer.data,
                status=status.HTTP_201_CREATED)
        return Response(
            serializer.errors,
            status=status.HTTP_400_BAD_REQUEST)
```

BROWSEABLE API

Snippet List

OPTIONS

GET ▾

This viewset automatically provides `list`, `create`, `retrieve`, `update` and `destroy` actions.

Additionally we also provide an extra `highlight` action.

`GET /snippets/`

HTTP 200 OK
Allow: `GET, POST, HEAD, OPTIONS`
Content-Type: `application/json`
Vary: `Accept`

```
{  
    "count": 1,  
    "next": null,  
    "previous": null,  
    "results": [  
        {  
            "url": "https://restframework.herokuapp.com/snippets/1/",  
            "id": 1,  
            "highlight": "https://restframework.herokuapp.com/snippets/1/highlight/",  
            "owner": "admin",  
            "title": "serializers.py",  
            "code": "from django.contrib.auth.models import User, Group\r\nfrom rest_framework import serializers\r\n\r\nclass UserSerializer(serializers.",  
            "linenos": false,  
            "language": "python",  
            "style": "default"  
        }  
    ]  
}
```

WITCHCRAFT

django-rest-witchcraft is an extension for Django REST Framework that adds support for SQLAlchemy. It aims to provide a similar development experience to building REST api's with Django REST Framework with Django ORM, except with SQLAlchemy.

— *Django REST Witchcraft*

SWAGGER

The OpenAPI Specification ... is the world's standard for defining RESTful interfaces. The OAS enables developers to design a technology-agnostic API interface that forms the basis of their API development and consumption.

Swagger.io

SWAGGER

Editor: Specify API via yaml

The screenshot shows the Swagger Editor interface split into two main sections. On the left, the **Swagger Editor** displays the YAML specification for the Petstore API. The code editor contains the following YAML:

```
1 swagger: "2.0"
2   info:
3     description: "This is a sample server Petstore server. You can find
4       out more about      Swagger at [http://swagger.io](http://swagger.io)
5       or on [irc.freenode.net, #swagger](http://swagger.io/irc/).      For
6       this sample, you can use the api key `special-key` to test the
7       authorization      filters."
8     version: "1.0.0"
9     title: "Swagger Petstore"
10    termsOfService: "http://swagger.io/terms/"
11    contact:
12      email: "apiteam@swagger.io"
13    license:
14      name: "Apache 2.0"
15      url: "http://www.apache.org/licenses/LICENSE-2.0.html"
16    host: "petstore.swagger.io"
17    basePath: "/v2"
18    tags:
19      - name: "pet"
20        description: "Everything about your Pets"
21        externalDocs:
22          description: "Find out more"
23          url: "http://swagger.io"
24      - name: "store"
25        description: "Access to Petstore orders"
26      - name: "user"
27        description: "Operations about user"
28        externalDocs:
29          description: "Find out more about our store"
30          url: "http://swagger.io"
31    schemes:
32      - "http"
33    paths:
34      /pet:
35        post:
36          tags:
37            - "pet"
38          summary: "Add a new pet to the store"
39          description: ""
```

On the right, the **Swagger Petstore** documentation is displayed. It includes a brief description of the Petstore server, links to the Terms of service, Contact the developer, and Apache 2.0, and a link to Find out more about Swagger. Below this, there's a section for Schemes with a dropdown set to HTTP and an Authorize button with a lock icon. The main content area shows the **pet** resource with its description: "Everything about your Pets". It lists two operations: a green **POST** operation for adding a new pet to the store, and an orange **PUT** operation for updating an existing pet.

SWAGGER

UI: Browseable API

The screenshot shows the Swagger UI interface for a RESTful API. At the top, there is a dropdown for 'Schemes' set to 'HTTP' and a 'Authorize' button with a lock icon. Below this, the 'pet' category is expanded, showing the following endpoints:

- POST /pet** Add a new pet to the store (green background)
- PUT /pet** Update an existing pet (orange background)
- GET /pet/findByStatus** Finds Pets by status (blue background)
- GET /pet/findByTags** Finds Pets by tags (light gray background)
- GET /pet/{petId}** Find pet by ID (blue background)
- POST /pet/{petId}** Updates a pet in the store with form data (green background)
- DELETE /pet/{petId}** Deletes a pet (red background)
- POST /pet/{petId}/uploadImage** uploads an image (green background)

Below the 'pet' section, the 'store' category is collapsed, indicated by a downward arrow.

SWAGGER

Codegen: Generate client sdk's automatically! (then test with them!)

```
12 Available Clients: [ akka-scala,
11   android, async-scala,clojure,cpprest,csharp,CsharpDotNet2,
10   cwiki,dart,dynamic-html,flash,go,groovy,html,
9   html2,java,javascript,javascript-closure-angular,
8   jaxrs-cxf-client,jmeter,objc,perl,php,python,
7   qt5cpp,ruby,scala,swagger,swagger-yaml,swift,
6   swift3,tizen,typescript-angular,typescript-angular2,
5   typescript-fetch,typescript-node],
4
3 Available Servers: [ aspnet5,aspnetcore,
2   erlang-server,go-server,haskell,inflector,
1   jaxrs,jaxrs-cxf,jaxrs-cxf-cdi,jaxrs-resteasy,
13   "jaxrs-spec","lumen","msf4j","nancyfx","nodejs-server",
1   python-flask,rails5,scalatra,silex-PHP, sinatra,
2   slim,spring,undertow]
~
```

SWAGGER

The API is data, whether generated from your backend or vice versa.
Metacode for client libraries provides consistency and removes boiler plate.

SWAGGER

Enables new workflows:

- DRF -> generate API spec -> generate client
- Write spec for existing API -> generate client
- Generate server stubs from spec, implement them

READING

- Python Metaclasses
- Documenting your API
- Browse: [Swagger API Design](#)